

Lecture 12:

Dynamic Memory Allocation 3

Rohan Murty

October 13, 2009




Road map for Today

- ▶ Mark and Sweep GC
- ▶ Generational Garbage Collection
- ▶ Reference Counting
- ▶ glibc malloc implementation
- ▶ Buddy allocation
- ▶ Measuring malloc performance

Mark and Sweep GC

- ▶ Idea: Use a **mark bit** in the header of each block
- ▶ GC scans all memory objects (starting at roots), and sets mark bit on each reachable memory block.
- ▶ Sweeping:
 - ▶ Scan over all memory blocks.
 - ▶ Any block without the mark bit set is freed




Reflecting on Mark and Sweep

► Advantages

- Easily handles cycles
- Low overhead

► Large object set

- Scanning all objects is *slow*
- **Expensive!**



Most objects have short lifespans

Most objects created post previous GC Cycle

*An object that survives a few GC cycles → Likely to survive longer*₄


Generational Garbage Collection

- ▶ Used in Java and .NET
- ▶ Fast incremental GC
 - ▶ Based on mark and sweep
 - ▶ Classify objects into 'generations'
 - ▶ GC certain generations more aggressively than others
- ▶ Scan “newer” objects more often than “older” ones

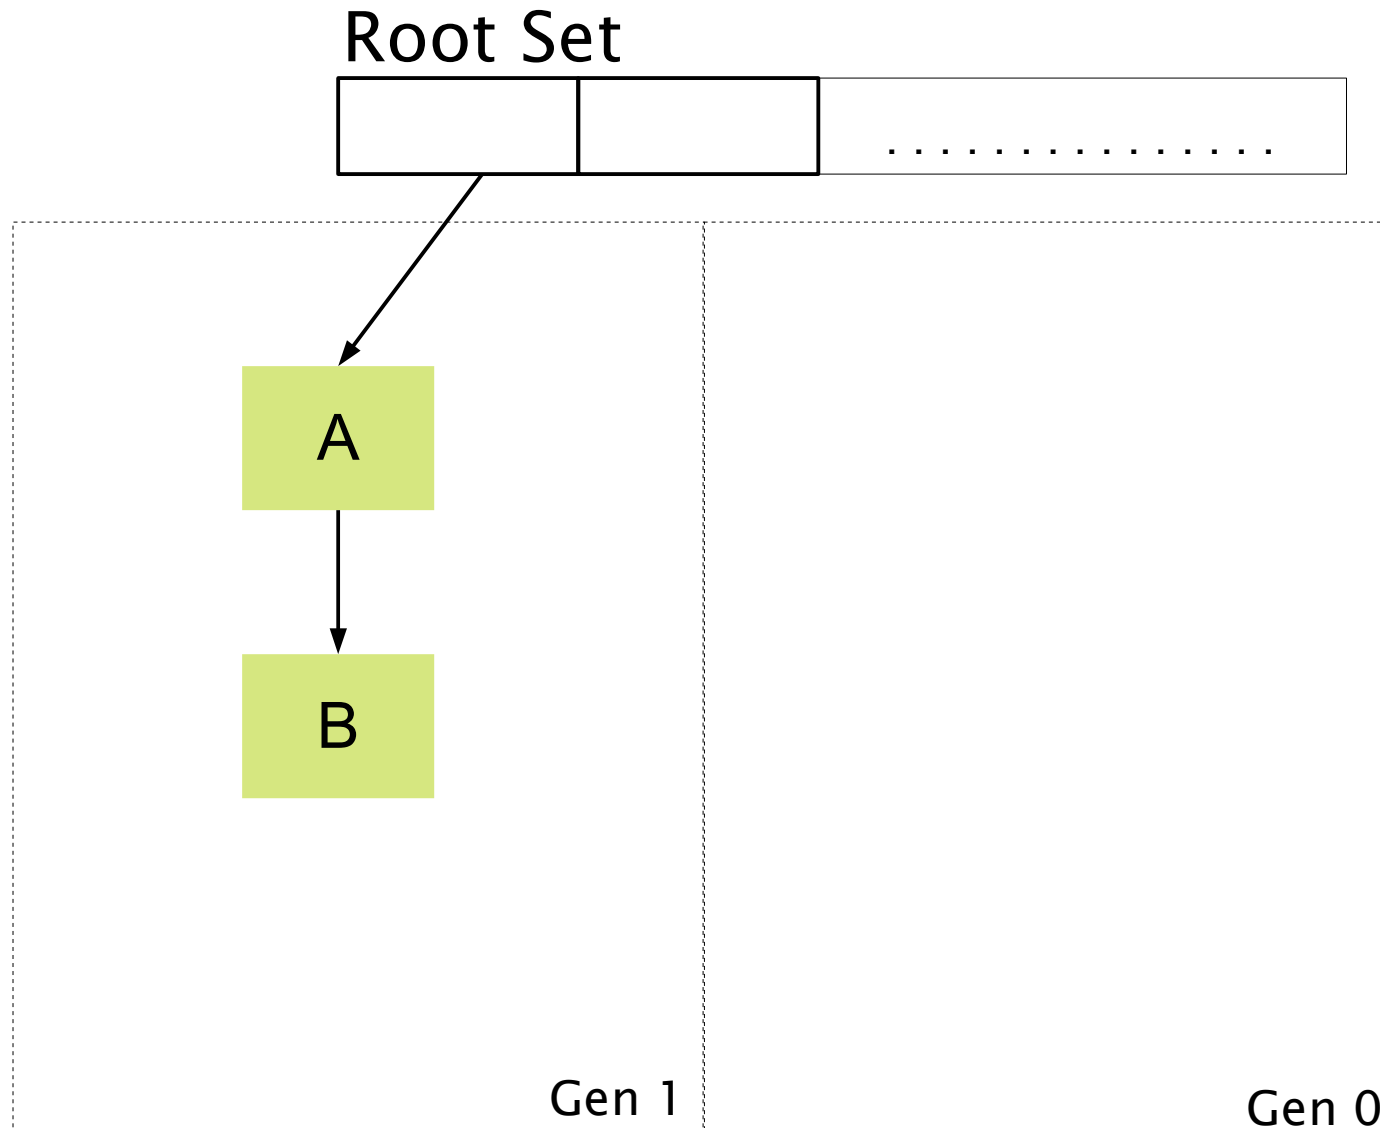
Generational Garbage Collection

- ▶ Algorithm:
- ▶ Divide objects into generations (Gen0, Gen1...)
- ▶ Aggressively GC lower generation objects
- ▶ If a region is filled with objects
 - ▶ Perform GC
 - ▶ For Objects of Gen(n) that survive GC
 - Promote to Gen(n+1)

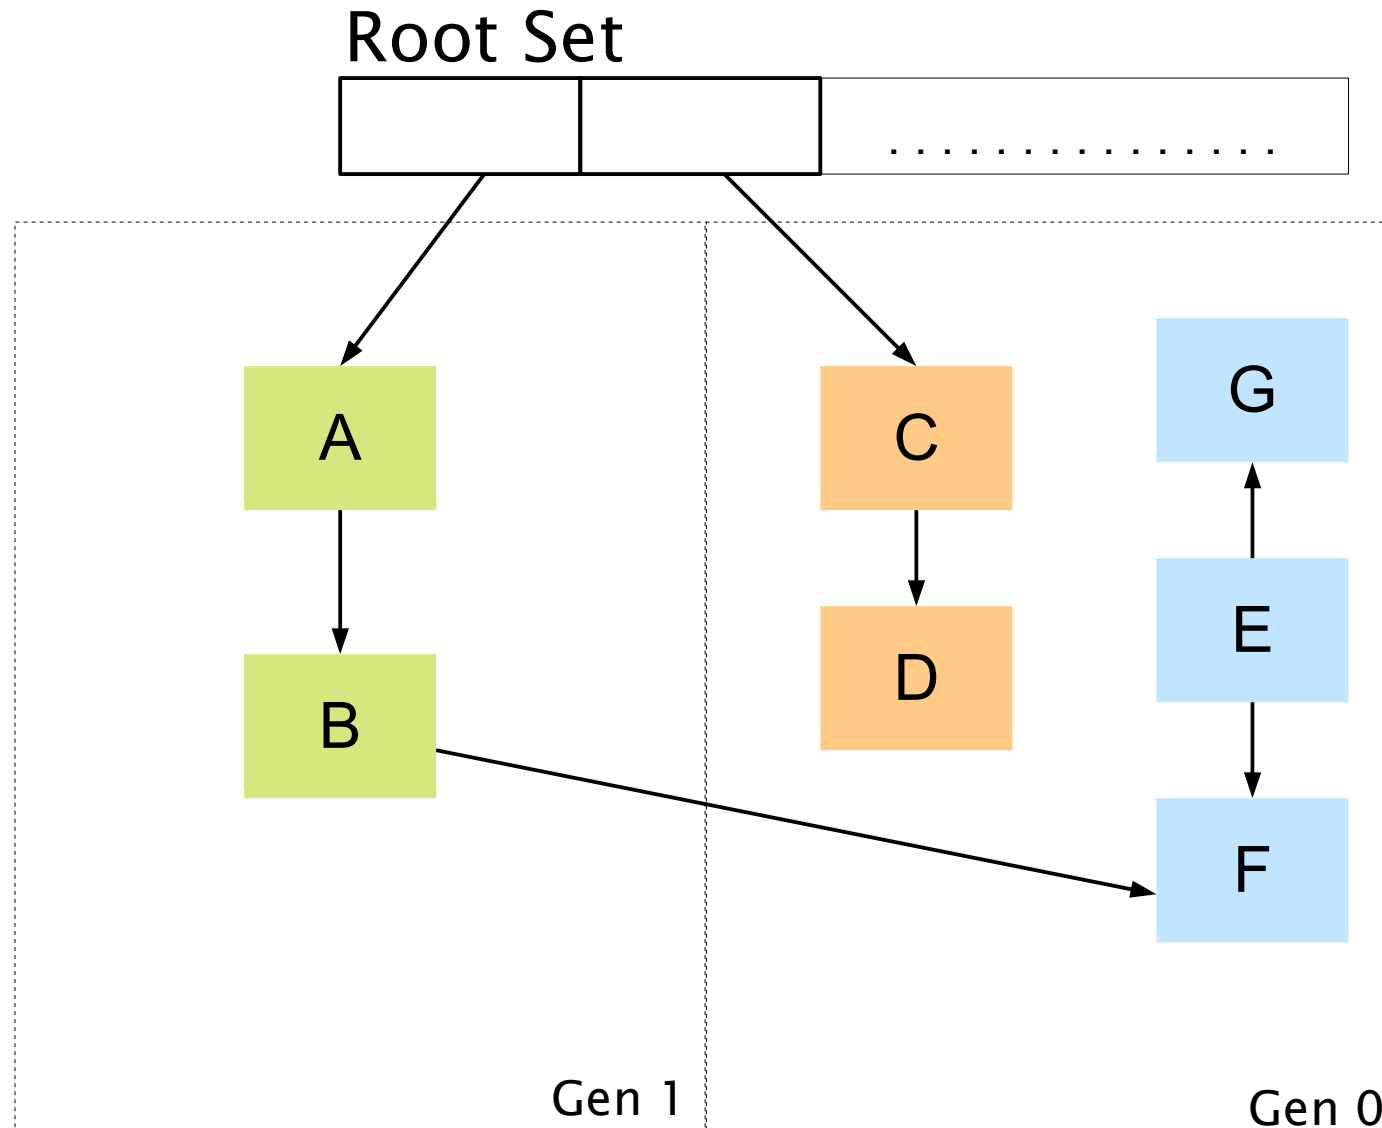
Generational Garbage Collection



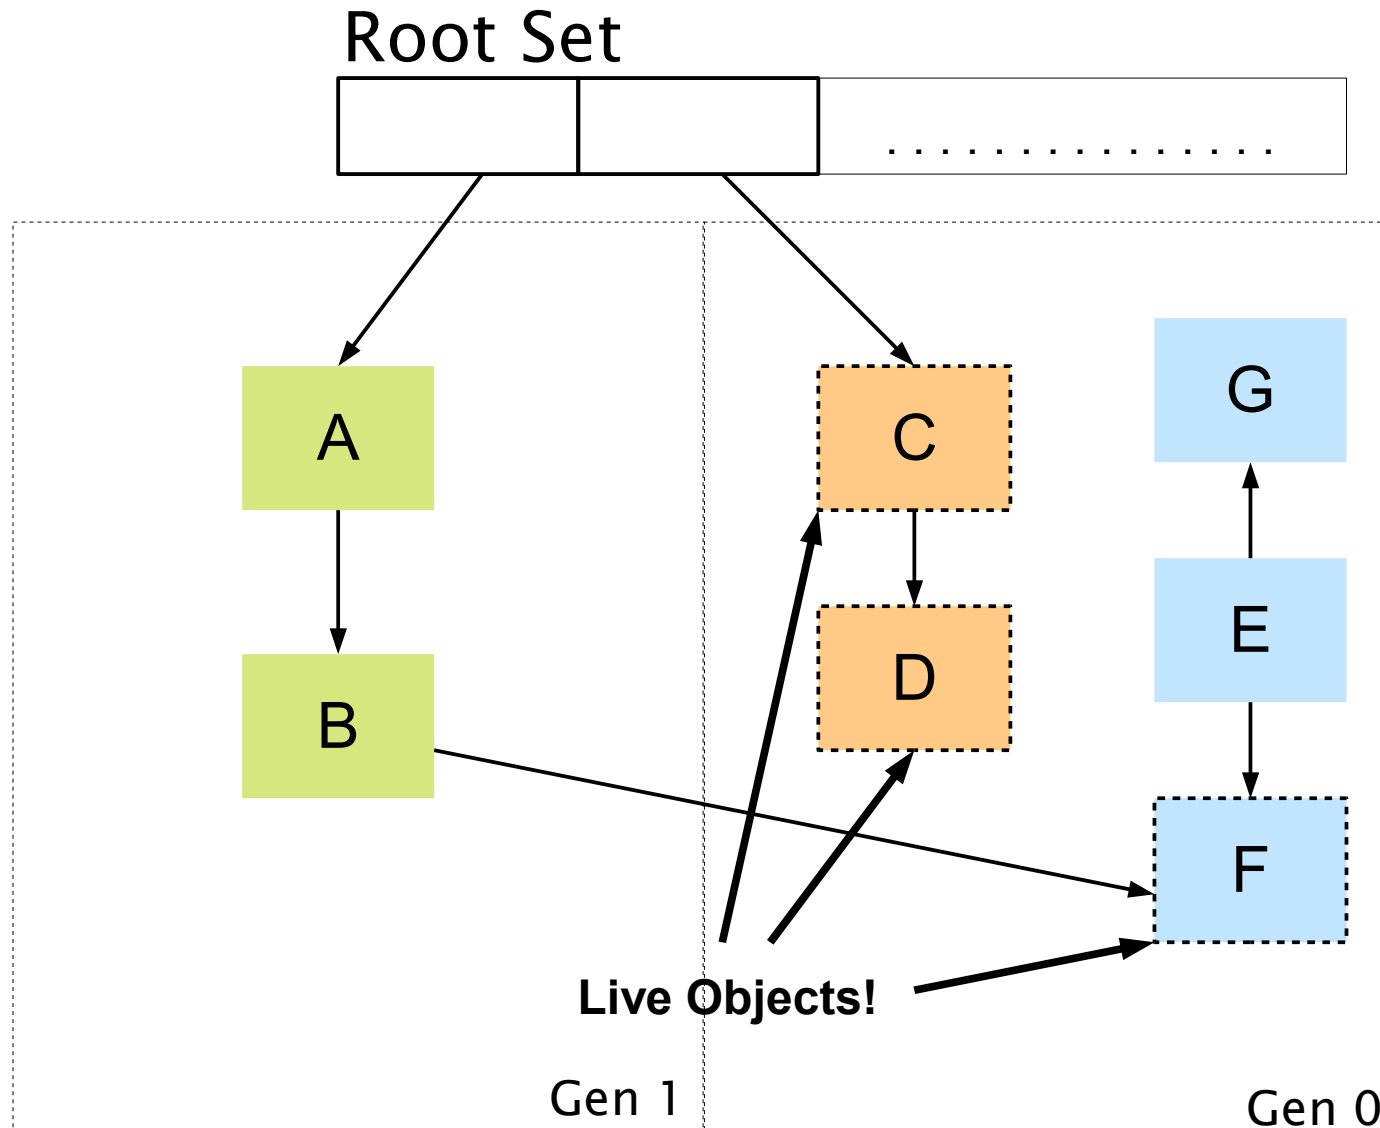
Generational Garbage Collection



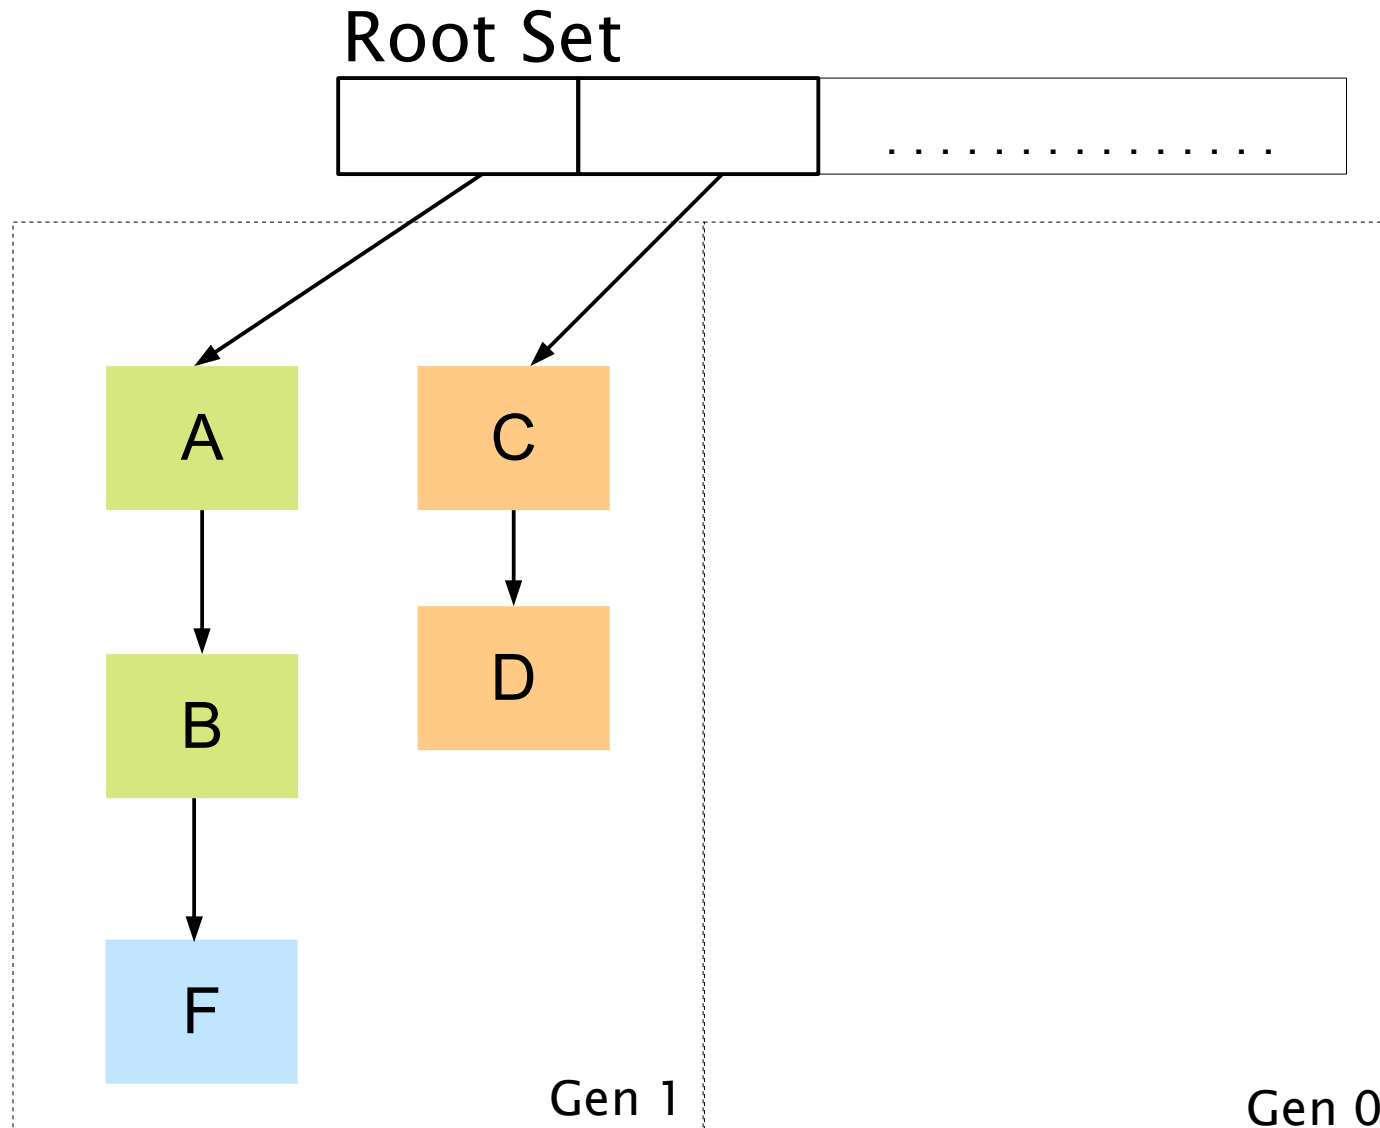
Generational Garbage Collection



Generational Garbage Collection



Generational Garbage Collection



Reference Counting

► **Problem:** How do we decide if an object is no longer referenced?

Reference Counting

- ▶ **Problem**: How do we decide if an object is no longer referenced?
- ▶ **Solution**: Keep track of number of *active* references to every object
- ▶ As soon as reference hits 0, can immediately GC object
- ▶ Tricky: How do we know an object is garbage, really?
- ▶ Used in: Python, Microsoft COM, Apple Cocoa, etc.
- ▶ Many C hash table implementations use ref. counting to keep track of active entries

Summarizing GC

- ▶ Root set used to track live objects
- ▶ Possible strategies
 - ▶ Mark and sweep
 - ▶ Generational GC
 - ▶ Reference Count
- ▶ C, by default, does not support GC
- ▶ Now, real world implementations of malloc

GNU C Library (glibc) malloc

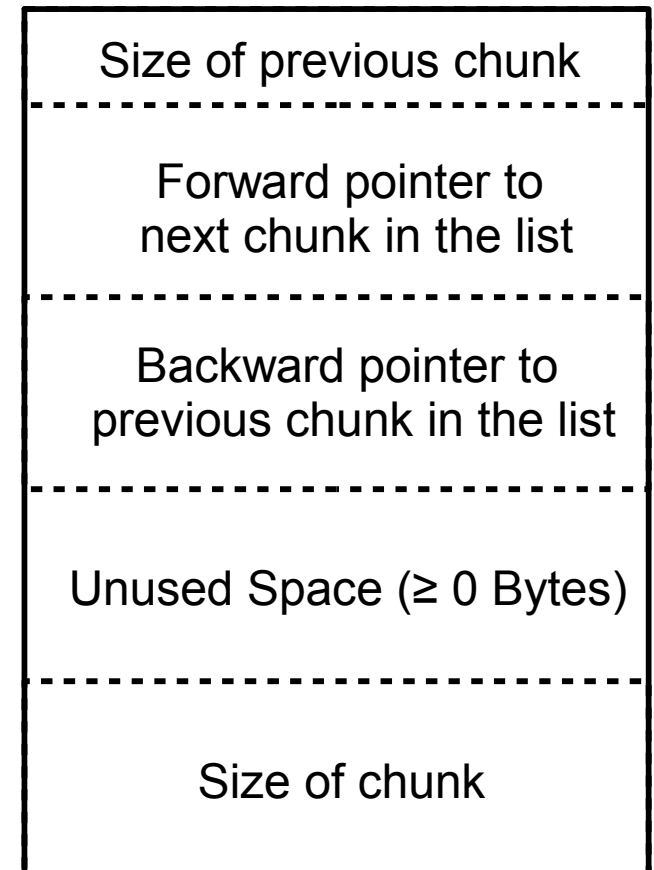
- ▶ Original implementation (dlmalloc)
- ▶ Current implementation (ptmalloc2)
 - ▶ Based on dlmalloc
 - ▶ Used in glibc 2.4 for linux based systems
 - ▶ Supports multiple heaps
 - ▶ But we won't worry about that for now

Doug Lea Allocator (dlmalloc)

- ▶ Fast and efficient implementation
- ▶ Heap allocated using `sbrk()`
- ▶ Memory allocated as “chunks”
- ▶ Uses a **best-fit** strategy
- ▶ Coalesce chunks during free
 - ▶ Reduces fragmentation
- ▶ Segregated lists (**binning**) to find free chunks quickly

Chunks

- ▶ Stored in bin
 - ▶ doubly-linked list
- ▶ Easy to coalesce adjacent chunks
 - ▶ Size of current and previous chunks stored in chunk header
- ▶ Minimum chunk size is 16 bytes




Free Chunk Format

Designing *dmalloc*

► Goal: Find **best-fitting** *free* chunk list

Designing *dlmalloc*

► Goal: Find **best-fitting** *free* chunk list




Designing *dldmalloc*

- ▶ Goal: Find **best-fitting** *free* chunk list
- ▶ *Observation: We receive many small (< 256 bytes) requests*

Segregated lists


Multiple bins



Key question: How do we decide bin-sizes?


Designing *dlmalloc*

- ▶ Goal: Find **best-fitting** *free* chunk list
- ▶ *Observation: We receive many small (< 256 bytes) requests*




Designing *dlmalloc*

Malloc (16 Bytes)




Designing *dlmalloc*



dlmalloc: Servicing Large Requests

► Search variable-sized bins

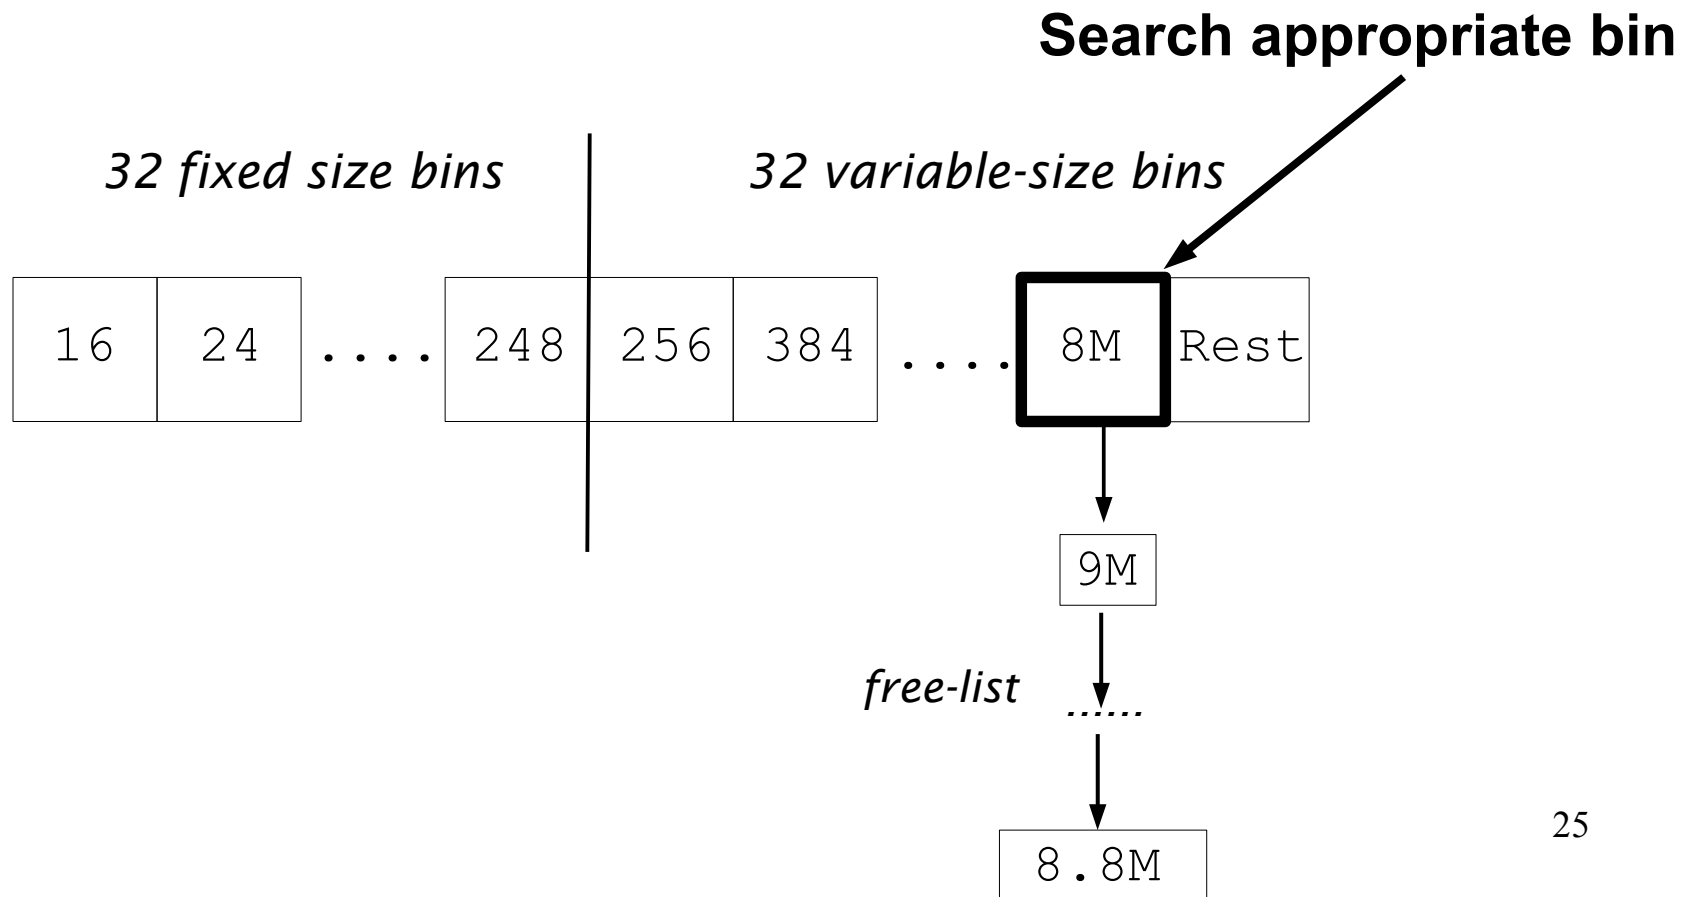
malloc (8.7MB)



dlmalloc: Servicing Large Requests

► Search variable-sized bins

malloc (8.7MB)



dlmalloc: Servicing Large Requests

► Search variable-sized bins

malloc (8.7MB)



dlmalloc: Optimizing for Large Requests

- ▶ For a given bin,
 - ▶ Linear search to find chunk size that fits
 - ▶ Potentially **expensive**

*We want to find **best-fit** chunk!*

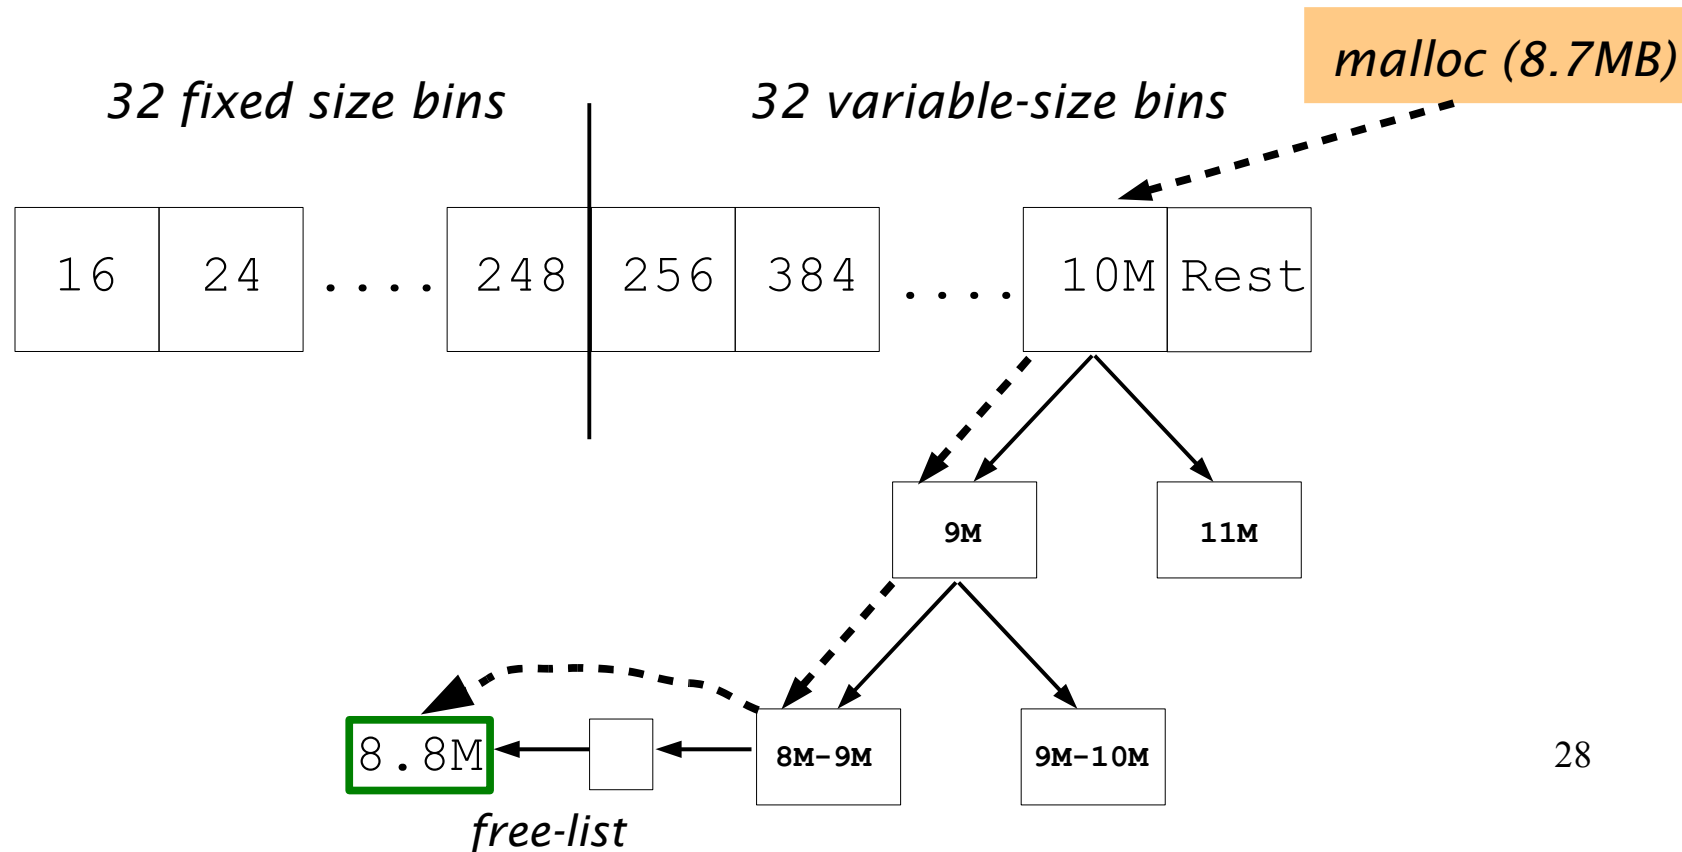
- ▶ Can we optimize the search for a chunk?
 - ▶ Sort free-list based on chunk sizes?
 - ▶ Sorting takes time
 - ▶ Impacts performance
- ▶ Can we optimize this further?

dlmalloc: Optimizing for Large Requests

- ▶ Search variable-sized bins
- ▶ Stores chunks as binary trees

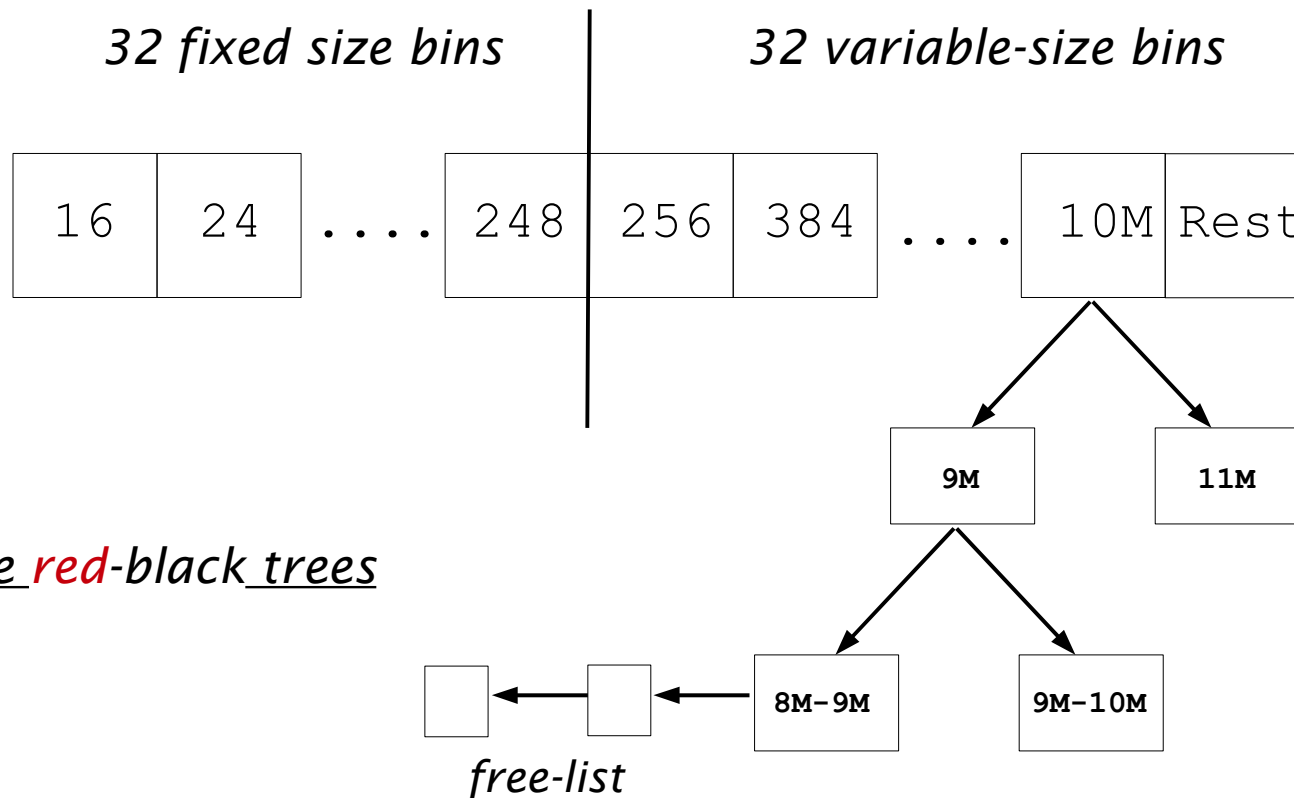
Avg. Search: $O(\log n)$

Avg. Insertion: $O(\log n)$



dlmalloc: Optimizing for Large Requests

- ▶ Search variable-sized bins
 - ▶ Stores chunks as binary trees
 - ▶ Each node stores list of chunks of *same* size



dlmalloc: Freeing Chunks

- ▶ Coalesce chunk with surrounding chunks
- ▶ Re-insert coalesced chunk into appropriate free list
- ▶ Fast coalescing because of 'boundary tags'
 - ▶ Size of chunk at the beginning
 - ▶ Size of chunk at the end

dlmalloc: Optimizing Performance

- ▶ Introducing quicklists/fastbins (< 64 bytes)
 - ▶ Fixed size bins for very small requests
- ▶ Treat different size chunks differently
 - ▶ Small chunks: do not coalesce, return to free list
 - ▶ Medium chunks: occasionally formed by coalescing objects from quicklists (optimistic strategy)
 - ▶ When free()-ed, performs immediate coalescing and occasionally splits to
 - Find best-fit
 - Allocate more space to quick lists
 - ▶ Large chunks: Coalescing is sometimes deferred

Summarizing dlmalloc

- ▶ Bins to store free lists (of chunks)
- ▶ Divide bins into two sets
 - ▶ First set, for small requests (occur often)
 - ▶ Second set, for larger requests (spaced at log. intervals)
 - ▶ Binary tree used to quickly insert and search chunks in second set
- ▶ Chunks store boundary tags
 - ▶ Coalesced during free
- ▶ Successive refinements performed over the years
 - ▶ Foundation for various malloc implementations
 - ▶ ptmalloc2, a variant of dlmalloc used in glibc 2.4

But how is malloc implemented for the rest of the world (80%+)?

Windows XP Allocator

- ▶ Best-fit search
- ▶ 127 exact-size quicklists
 - ▶ For multiple of 8 bytes
- ▶ Objects > 1024 bytes
 - ▶ Obtained from a sorted list
 - ▶ Sacrifice speed for a good fit



Buddy (Memory) Allocation

- ▶ Described by Don Knuth
 - ▶ First proposed by economist Harry Markowitz
- ▶ Very fast and simple for 2^n size blocks
- ▶ **Idea:** *split* memory into *halves* until **best-fit** is found
 - ▶ When **free**-ing
 - ▶ Combine adjacent empty halves

Buddy (Memory) Allocation: Example

Lower bound for block size 5 $\rightarrow 2^4 = 16K$ (*minimum* block size)

Upper bound for block size 9 $\rightarrow 2^9 = 512K$ (**maximum** block size)


	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	384K											

malloc (20K)

Simple Strategy: Allocate the **entire** 384K \rightarrow results in *severe* internal fragmentation

Idea: ***Split*** memory until we get *smallest* (permissible) block that satisfies the request

Buddy Allocation: Splitting Memory



.....
.....
.....

Buddy (Memory) Allocation: Example

	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	32K	16K	48K	96K			192K					

Buddy (Memory) Allocation: Example

	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	32K	16K	48K	96K				192K				

malloc (38K)

No need to split any blocks. 48K block will suffice!

Buddy (Memory) Allocation: Example

	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	32K	16K	48K	96K			192K					

malloc (89K)

Buddy (Memory) Allocation: Example

	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	32K	16K	48K	96K			192K					

malloc (15K)

Buddy (Memory) Allocation: Example

	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	32K	16K	48K	96K			192K					


free

Buddy (Memory) Allocation: Example

	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K	32K
Memory	32K	16K	48K	96K			192K					


free

Buddy (Memory) Allocation: Example



Opportunity to coalesce to adjacent free blocks (32K and 16K)

Buddy (Memory) Allocation: Example



Buddy Memory Allocation

► Pros:

- Easy to implement (binary tree)
- Little external fragmentation
 - Why??

► Cons:

- Internal fragmentation
 - Why??

Not all requests will be on the order of $2^k \rightarrow$ More than requested memory allocated

Bounds on block sizes impact internal fragmentation

What constitutes a *good* malloc impl.?

▶ Time

- ▶ Minimize time taken for allocation and free

▶ Space

- ▶ Minimize fragmentation

▶ Locality

- ▶ Allocations around time t are stored near each other
- ▶ Similar sized blocks are stored near by

Performance: Throughput

- ▶ Given a sequence of requests

malloc(...) , malloc (...), free (...), malloc(...)

- ▶ Measure: Throughput

- ▶ Number of requests satisfied per unit time

- ▶ Goal: Maximize throughput

Performance: Peak Utilization

- ▶ Given a sequence of requests

malloc(...) , malloc (...), free (...), malloc(...)

- ▶ Measure: Peak Utilization

- ▶ Maximum currently allocated memory / Current Heap Size

- ▶ Goal: Maximize peak utilization

Overall Performance Goal

- ▶ Maximize throughput *and* peak utilization
- ▶ Conflict with each other. Why?
 - ▶ Throughput: Minimize time taken to allocate
 - ▶ Peak Utilization: Minimize waste of space
 - ▶ Need more time to do this efficiently
 - ▶ Conflicts with the Throughput goal
- ▶ Find sweet spot between
 - ▶ maximizing throughput and minimizing peak utilization

SPEC Benchmarks

- ▶ Popular tool used to benchmark system performance
 - ▶ Returns a single performance metric: *time*
- ▶ Suite of applications to stress test a system
 - ▶ Mesh generation, Monte carlo, circuit simulation and analysis, fluid dynamics applications, compression, neural networks, etc.
- ▶ Manipulates large and small arrays in loops
 - ▶ Sensitive to performance of memory allocation
 - ▶ Good stress test for malloc

References

- ▶ Lecture notes: *Sebastian Hack* and *Christoph Weidenbach (MPI)*
- ▶ “Understanding the heap by breaking it” - *Justin Ferguson*
- ▶ “Back to basics: Generational GC” - *Abhinaba, Microsoft Corp*
- ▶ “A Locality Improving Dynamic Memory Allocator” - *Yi Eng and Emery Berger, MSP '05*
- ▶ “Reconsidering Custom Memory Allocation” - *Berger, Zorn, and McKinley, OOPSLA 2002*
- ▶ Wikipedia